

MSSC 6250 Machine Learning Homework 3

Bayesian Regression, Logistic Regression, Generative Models

Henri Medeiros Dos Reis

- Deadline: **Friday, April 7 11:59 PM**
- Homework presentation date: **Tuesday, April 11**
- Please submit your work in **one PDF** file to **D2L > Assessments > Dropbox**. *Multiple files or a file that is not in pdf format are not allowed.*
- Any relevant code should be attached.
- Read **ISL** Chapter 4.

Exercises required for all students

1. ISL Sec. 4.8: 2

Solution:

4.17 since all σ are equal gives

$$p_k(x) = \frac{\pi_k \frac{1}{\sqrt{2\pi}\sigma} \exp(-\frac{1}{2\sigma^2}(x - \mu_k)^2)}{\sum_{i=1}^K \pi_i \frac{1}{\sqrt{2\pi}\sigma} \exp(-\frac{1}{2\sigma^2}(x - \mu_i)^2)}$$

and 4.18

$$\delta_k(x) = x \frac{\mu_k}{\sigma^2} - \frac{\mu_k^2}{2\sigma^2} + \log(\pi_k)$$

Given $p_k(x)$ is normal. We need to classify an observation to class k , maximizing $p_k(x)$. We can take the log, since it is equivalent to the function when maximizing

$$\log(p_k(x)) = \log \left(\frac{\pi_k \exp(-\frac{1}{2\sigma^2}(x - \mu_k)^2)}{\sum_{i=1}^K \pi_i \exp(-\frac{1}{2\sigma^2}(x - \mu_i)^2)} \right) =$$

$$\begin{aligned} & \log(\pi_k \exp(-\frac{1}{2\sigma^2}(x - \mu_k)^2)) - \log(\sum_{i=1}^K \pi_i \exp(-\frac{1}{2\sigma^2}(x - \mu_i)^2)) = \\ & \log(\pi_k) + (-\frac{1}{2\sigma^2}(x - \mu_k)^2) - \log(\sum_{i=1}^K \pi_i \exp(-\frac{1}{2\sigma^2}(x - \mu_i)^2)) \end{aligned}$$

Since the third term is just a constant, it would disappear when we take the derivative with respect to k , since it is a sum over K . Therefore, maximizing the first two terms is equivalent to maximizing the whole equation. Then

$$\begin{aligned} \Rightarrow \log(\pi_k) + (-\frac{1}{2\sigma^2}(x - \mu_k)^2) &= \log(\pi_k) - \frac{x^2 - 2\mu_k x + \mu_k^2}{2\sigma^2} \\ &= \log(\pi_k) - \frac{\mu_k^2}{2\sigma^2} + \frac{\mu_k x}{\sigma^2} - \frac{x^2}{2\sigma^2} \end{aligned}$$

And $\frac{x^2}{2\sigma^2}$ is a constant

$$\Rightarrow \log(\pi_k) - \frac{\mu_k^2}{2\sigma^2} + \frac{\mu_k x}{\sigma^2} = x \frac{\mu_k}{\sigma^2} - \frac{\mu_k^2}{2\sigma^2} + \log(\pi_k)$$

Which is exactly the same as 4.18. Therefore, maximizing $p_k(x)$ is the same as maximizing $\delta_k(x)$

2. ISL Sec. 4.8: 3

Solution:

Now we follow the same steps as last problem, but with different σ s

$$\begin{aligned} p_k(x) &= \frac{\pi_k \frac{1}{\sqrt{2\pi}\sigma_k} \exp(-\frac{1}{2\sigma_k^2}(x - \mu_k)^2)}{\sum_{i=1}^K \pi_i \frac{1}{\sqrt{2\pi}\sigma_i} \exp(-\frac{1}{2\sigma_i^2}(x - \mu_i)^2)} \\ &= \frac{\frac{\pi_k}{\sigma_k} \exp(-\frac{1}{2\sigma_k^2})(x - \mu_k)^2}{\sum_{l=1}^K \frac{\pi_l}{\sigma_l} \exp(-\frac{1}{2\sigma_l^2})(x - \mu_l)^2} \\ \Rightarrow \log(p_k(x)) &= \log \left(\frac{\frac{\pi_k}{\sigma_k} \exp(-\frac{1}{2\sigma_k^2})(x - \mu_k)^2}{\sum_{l=1}^K \frac{\pi_l}{\sigma_l} \exp(-\frac{1}{2\sigma_l^2})(x - \mu_l)^2} \right) \\ &= \log(\frac{\pi_k}{\sigma_k} \exp(-\frac{1}{2\sigma_k^2})(x - \mu_k)^2) - \log(\sum_{l=1}^K \frac{\pi_l}{\sigma_l} \exp(-\frac{1}{2\sigma_l^2})(x - \mu_l)^2) \end{aligned}$$

and the term involving σ_l is a constant

$$\Rightarrow \log(\pi_k) - \frac{1}{2\sigma_k^2}(x - \mu_k^2) - \log(\sigma_k)$$

$$= \log(\pi_k) - \frac{\mu_k^2}{2\sigma_k^2} + \frac{\mu_k x}{\sigma_k^2} - \frac{x^2}{2\sigma_k^2} \log(\sigma_k)$$

As we can see, we have terms involving x that are squared. Therefore, not linear, and in fact, it is quadratic

3. ISL Sec. 4.8: 5

Solution:

a-) QDA would perform better on the training data set, because of its flexibility. But the opposite should happen in the test data set, LDA would be better, because QDA would overfit.

b-) QDA would perform better in both training and testing data sets in most cases. Since the flexibility would capture the non-linearity better in most cases. However it highly depends on the nature of the non-linearity.

c-) The test predict accuracy of QDA would improve compared to LDA because of the flexibility advantage of QDA.

d-) False, with a small enough sample size, the QDA model will easily overfit, leading to a larger test error compared to LDA.

4. ISL Sec. 4.8: 7

Solution:

$$f_{yes}(x) = \frac{e^{\frac{(x-\mu)^2}{2\sigma^2}}}{\sqrt{2\pi\sigma^2}}$$

$$f_{yes}(4) = \frac{e^{\frac{-36}{72}}}{\sqrt{72\pi}}$$

$$f_{no}(x) = \frac{e^{\frac{(x-\mu)^2}{2\sigma^2}}}{\sqrt{2\pi\sigma^2}}$$

$$f_{no}(4) = \frac{e^{\frac{-2}{9}}}{\sqrt{72\pi}}$$

$$P_r(Y = yes|X = 4) = \frac{\pi_{yes}f_{yes}(4)}{\pi_{no}f_{no}(4) + \pi_{yes}f_{yes}(4)}$$

$$= \frac{0.8 \frac{e^{\frac{-36}{72}}}{\sqrt{72\pi}}}{0.2 \frac{e^{\frac{-2}{9}}}{\sqrt{72\pi}} + 0.8 \frac{e^{\frac{-36}{72}}}{\sqrt{72\pi}}}$$

5. ISL Sec. 4.8: 10

Solution:

$$\begin{aligned}
 \log\left(\frac{P_r(Y = k|X = x)}{P_r(Y = K|X = x)}\right) &= \log\left(\frac{\pi_k f_k(x)}{\pi_K f_K(x)}\right) \\
 &= \log\left(\frac{\pi_k}{\pi_K}\right) + \log\left(\frac{\frac{1}{\sqrt{2\pi}\sigma} \exp(\frac{1}{2\sigma^2}(x - \mu_k)^2)}{\frac{1}{\sqrt{2\pi}\sigma} \exp(\frac{1}{2\sigma^2}(x - \mu_K)^2)}\right) \\
 &= \log\left(\frac{\pi_k}{\pi_K}\right) - \frac{(x - \mu_k)^2}{2\sigma^2} + \frac{(x - \mu_K)^2}{2\sigma^2} \\
 &= \log\left(\frac{\pi_k}{\pi_K}\right) - \frac{x^2}{2\sigma^2} + \frac{\mu_k x}{\sigma^2} - \frac{\mu_k^2}{2\sigma^2} + \frac{x^2}{2\sigma^2} - \frac{\mu_K x}{\sigma^2} + \frac{\mu_K^2}{2\sigma^2} \\
 &= \log\left(\frac{\pi_k}{\pi_K}\right) - \frac{\mu_k^2}{2\sigma^2} + \frac{\mu_K^2}{2\sigma^2} + \frac{\mu_k x}{\sigma^2} - \frac{\mu_K x}{\sigma^2} \\
 &= \log\left(\frac{\pi_k}{\pi_K}\right) + \frac{\mu_K^2 - \mu_k^2}{2\sigma^2} + x\left(\frac{\mu_k - \mu_K}{\sigma^2}\right) \\
 &= a_k + x b_{kj}
 \end{aligned}$$

Where $a_k = \log\left(\frac{\pi_k}{\pi_K}\right) + \frac{\mu_K^2 - \mu_k^2}{2\sigma^2}$ and $b_{kj} = \frac{\mu_k - \mu_K}{\sigma^2}$

6. ISL Sec. 4.8: 13

Solution:

a-)

```
library(ISLR2)
```

Warning: package 'ISLR2' was built under R version 4.2.2

```
head(Weekly)
```

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	Today	Direction
1	1990	0.816	1.572	-3.936	-0.229	-3.484	0.1549760	-0.270	Down
2	1990	-0.270	0.816	1.572	-3.936	-0.229	0.1485740	-2.576	Down
3	1990	-2.576	-0.270	0.816	1.572	-3.936	0.1598375	3.514	Up
4	1990	3.514	-2.576	-0.270	0.816	1.572	0.1616300	0.712	Up
5	1990	0.712	3.514	-2.576	-0.270	0.816	0.1537280	1.178	Up
6	1990	1.178	0.712	3.514	-2.576	-0.270	0.1544440	-1.372	Down

```
cor(Weekly[, -9])
```

	Year	Lag1	Lag2	Lag3	Lag4
Year	1.00000000	-0.032289274	-0.03339001	-0.03000649	-0.031127923
Lag1	-0.03228927	1.000000000	-0.07485305	0.05863568	-0.071273876
Lag2	-0.03339001	-0.074853051	1.00000000	-0.07572091	0.058381535
Lag3	-0.03000649	0.058635682	-0.07572091	1.00000000	-0.075395865
Lag4	-0.03112792	-0.071273876	0.05838153	-0.07539587	1.000000000
Lag5	-0.03051910	-0.008183096	-0.07249948	0.06065717	-0.075675027
Volume	0.84194162	-0.064951313	-0.08551314	-0.06928771	-0.061074617
Today	-0.03245989	-0.075031842	0.05916672	-0.07124364	-0.007825873

	Lag5	Volume	Today
Year	-0.030519101	0.84194162	-0.032459894
Lag1	-0.008183096	-0.06495131	-0.075031842
Lag2	-0.072499482	-0.08551314	0.059166717
Lag3	0.060657175	-0.06928771	-0.071243639
Lag4	-0.075675027	-0.06107462	-0.007825873
Lag5	1.000000000	-0.05851741	0.011012698
Volume	-0.058517414	1.00000000	-0.033077783
Today	0.011012698	-0.03307778	1.000000000

summary(Weekly)

Year		Lag1		Lag2		Lag3	
Min.	:1990	Min.	:-18.1950	Min.	:-18.1950	Min.	:-18.1950
1st Qu.:	:1995	1st Qu.:	-1.1540	1st Qu.:	-1.1540	1st Qu.:	-1.1580
Median	:2000	Median	: 0.2410	Median	: 0.2410	Median	: 0.2410
Mean	:2000	Mean	: 0.1506	Mean	: 0.1511	Mean	: 0.1472
3rd Qu.:	:2005	3rd Qu.:	1.4050	3rd Qu.:	1.4090	3rd Qu.:	1.4090
Max.	:2010	Max.	: 12.0260	Max.	: 12.0260	Max.	: 12.0260

Lag4		Lag5		Volume		Today	
Min.	:-18.1950	Min.	:-18.1950	Min.	:0.08747	Min.	:-18.1950
1st Qu.:	-1.1580	1st Qu.:	-1.1660	1st Qu.:	0.33202	1st Qu.:	-1.1540
Median	: 0.2380	Median	: 0.2340	Median	:1.00268	Median	: 0.2410
Mean	: 0.1458	Mean	: 0.1399	Mean	:1.57462	Mean	: 0.1499
3rd Qu.:	1.4090	3rd Qu.:	1.4050	3rd Qu.:	2.05373	3rd Qu.:	1.4050
Max.	: 12.0260	Max.	: 12.0260	Max.	:9.32821	Max.	: 12.0260

Direction
Down:484
Up :605

Year and volume appear to be correlated. From the summary statistics, we can observe that all the Lag variables are somewhat similar to each other and 'Today'.

b-)

```
logit_fit = glm(Direction ~ Lag1+Lag2+Lag3+Lag4+Lag5+Volume,
                 data=Weekly, family=binomial)
summary(logit_fit)
```

Call:

```
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
     Volume, family = binomial, data = Weekly)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-1.6949	-1.2565	0.9913	1.0849	1.4579

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.26686	0.08593	3.106	0.0019 **
Lag1	-0.04127	0.02641	-1.563	0.1181
Lag2	0.05844	0.02686	2.175	0.0296 *
Lag3	-0.01606	0.02666	-0.602	0.5469
Lag4	-0.02779	0.02646	-1.050	0.2937
Lag5	-0.01447	0.02638	-0.549	0.5833
Volume	-0.02274	0.03690	-0.616	0.5377

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1496.2 on 1088 degrees of freedom
Residual deviance: 1486.4 on 1082 degrees of freedom
AIC: 1500.4

Number of Fisher Scoring iterations: 4

Looking at the p values, only lag 2 appears to be significant.

c-)

```
pred_prob <- predict(logit_fit, type = "response")
(res<-table(pred_prob > 0.5, Weekly$Direction))
```

	Down	Up
FALSE	54	48
TRUE	430	557

```
correct <- res[1,1]+res[2,2]
total <- sum(res)
correct/total
```

```
[1] 0.5610652
```

Only 56% of the predictions was correct. Considering that a random guess would give around 50%, this model is not much better.

d-)

```
train <- Weekly[Weekly$Year <= 2008, ]
test <- Weekly[Weekly$Year > 2008, ]

logit_fit_lag2 <- glm(Direction ~ Lag2, data = train, family = "binomial")

pred_prob <- factor(ifelse(
  predict(logit_fit_lag2, newdata = test, type = "response") < 0.5, "Down", "Up"))

caret::confusionMatrix(pred_prob, test$Direction, positive = "Up")
```

Confusion Matrix and Statistics

	Reference	
Prediction	Down	Up
Down	9	5
Up	34	56

Accuracy : 0.625
 95% CI : (0.5247, 0.718)
 No Information Rate : 0.5865
 P-Value [Acc > NIR] : 0.2439

 Kappa : 0.1414

 McNemar's Test P-Value : 7.34e-06

 Sensitivity : 0.9180
 Specificity : 0.2093
 Pos Pred Value : 0.6222
 Neg Pred Value : 0.6429
 Prevalence : 0.5865
 Detection Rate : 0.5385
 Detection Prevalence : 0.8654

Balanced Accuracy : 0.5637

'Positive' Class : Up

This model is an improvement from the previous one, but not by much. The accuracy is at 62.5%.

e-)

```
lda_fit_lag2 <- MASS::lda(Direction ~ Lag2, data = train, family = "binomial")  
pred_prob_lda <- predict(lda_fit_lag2, newdata = test, type = "response")  
caret::confusionMatrix(data=pred_prob_lda$class, test$Direction, positive = "Up")
```

Confusion Matrix and Statistics

	Reference	
Prediction	Down	Up
Down	9	5
Up	34	56

Accuracy : 0.625

95% CI : (0.5247, 0.718)

No Information Rate : 0.5865

P-Value [Acc > NIR] : 0.2439

Kappa : 0.1414

Mcnemar's Test P-Value : 7.34e-06

Sensitivity : 0.9180

Specificity : 0.2093

Pos Pred Value : 0.6222

Neg Pred Value : 0.6429

Prevalence : 0.5865

Detection Rate : 0.5385

Detection Prevalence : 0.8654

Balanced Accuracy : 0.5637

'Positive' Class : Up

This model is an improvement from the first model, but give pretty much the same results as the previous one. The accuracy is at 62.5%.

f-)

```
qda_fit_lag2 <- MASS::qda(Direction ~ Lag2, data = train, family = "binomial")  
pred_prob_qda <- predict(qda_fit_lag2, newdata = test, type = "response")  
caret::confusionMatrix(data=pred_prob_qda$class, test$Direction, positive = "Up")
```

Confusion Matrix and Statistics

```
              Reference  
Prediction Down Up  
      Down      0  0  
      Up       43 61  
  
      Accuracy : 0.5865  
      95% CI : (0.4858, 0.6823)  
No Information Rate : 0.5865  
P-Value [Acc > NIR] : 0.5419  
  
      Kappa : 0  
  
McNemar's Test P-Value : 1.504e-10  
  
      Sensitivity : 1.0000  
      Specificity : 0.0000  
Pos Pred Value : 0.5865  
Neg Pred Value :      NaN  
Prevalence : 0.5865  
Detection Rate : 0.5865  
Detection Prevalence : 1.0000  
Balanced Accuracy : 0.5000  
  
'Positive' Class : Up
```

We get an Accuracy of 58.6%.

g-)

```
set.seed(1)  
predicted_knn <- class::knn(train = data.frame(Lag2 = train$Lag2),  
                           test = data.frame(Lag2 = test$Lag2),  
                           cl = train$Direction,  
                           k = 1,
```

```

        prob = T)

caret::confusionMatrix(data = predicted_knn,
                        reference = test$Direction,
                        positive = "Up")

```

Confusion Matrix and Statistics

	Reference	
Prediction	Down	Up
Down	21	30
Up	22	31

```

Accuracy : 0.5
95% CI : (0.4003, 0.5997)
No Information Rate : 0.5865
P-Value [Acc > NIR] : 0.9700

```

```

Kappa : -0.0033

```

```

Mcnemar's Test P-Value : 0.3317

```

```

Sensitivity : 0.5082
Specificity : 0.4884
Pos Pred Value : 0.5849
Neg Pred Value : 0.4118
Prevalence : 0.5865
Detection Rate : 0.2981
Detection Prevalence : 0.5096
Balanced Accuracy : 0.4983

```

```

'Positive' Class : Up

```

This model gave an accuracy of 50% which is basically the same as guessing.
h-)

```

nb_fit <- e1071::naiveBayes(Direction~Lag2 ,data=train)

nb_pred <- predict(nb_fit, test)

caret::confusionMatrix(data = nb_pred,
                        reference = test$Direction,

```

```
positive = "Up")
```

Confusion Matrix and Statistics

```
      Reference
Prediction Down Up
Down      0   0
Up       43  61
```

```
Accuracy : 0.5865
 95% CI : (0.4858, 0.6823)
No Information Rate : 0.5865
P-Value [Acc > NIR] : 0.5419
```

```
Kappa : 0
```

```
McNemar's Test P-Value : 1.504e-10
```

```
Sensitivity : 1.0000
Specificity : 0.0000
Pos Pred Value : 0.5865
Neg Pred Value : NaN
Prevalence : 0.5865
Detection Rate : 0.5865
Detection Prevalence : 1.0000
Balanced Accuracy : 0.5000
```

```
'Positive' Class : Up
```

We get an Accuracy of 58.6% which is the same as we got with the QDA.

i-)

LDA & Logistic Regression get the same test accuracy of 0.625, so these two are tied as the best model.

j-)

```
predicted_knn6 <- class::knn(train = data.frame(Lag2 = train$Lag2),
                             test = data.frame(Lag2 = test$Lag2),
                             cl = train$Direction,
                             k = 6,
                             prob = T)
```

```
cm <- caret::confusionMatrix(data = predicted_knn6,  
                             reference = test$Direction,  
                             positive = "Up")  
cm$overall[1]
```

Accuracy
0.5769231

Trying KNN with 6 nearest neighbors gives an accuracy of 56.7%

```
predicted_knn9 <- class::knn(train = data.frame(Lag2 = train$Lag2),  
                             test = data.frame(Lag2 = test$Lag2),  
                             cl = train$Direction,  
                             k = 9,  
                             prob = T)  
  
cm <- caret::confusionMatrix(data = predicted_knn9,  
                             reference = test$Direction,  
                             positive = "Up")  
cm$overall[1]
```

Accuracy
0.5576923

Trying KNN with 9 nearest neighbors gives an accuracy of 55.7%

```
predicted_knn15 <- class::knn(train = data.frame(Lag2 = train$Lag2),  
                              test = data.frame(Lag2 = test$Lag2),  
                              cl = train$Direction,  
                              k = 15,  
                              prob = T)  
  
cm <- caret::confusionMatrix(data = predicted_knn15,  
                              reference = test$Direction,  
                              positive = "Up")  
cm$overall[1]
```

Accuracy
0.5865385

Trying KNN with 15 nearest neighbors gives an accuracy of 58.7%

The results seem to get better as we increase the number of nearest neighbors. Which is likely going to stop at some point and start to get worse as it increases.

7. ISL Sec. 4.8: 14


Solution:

a-)

```
mpg01 = rep(0, length(Auto$mpg))
mpg01[Auto$mpg > median(Auto$mpg)] = 1
Auto = data.frame(Auto, mpg01)
```

b-)

```
pairs(Auto)
```



We can see that there is a positive correlation between mpg and mpg01, and a negative correlation between cylinders, displacement, weight, horsepower and mpg01.

c-)

```
index <- caret::createDataPartition(y = Auto$mpg01, p = 0.8, list = F)

train1 <- Auto[index, ]
test1 <- Auto[-index, ]
```

d-)

```
lda_fit_mpg <- MASS::lda(mpg01~cylinders+weight+displacement
                        +horsepower,
                        data = train1, family = "binomial")

pred_prob_lda_mpg <- predict(lda_fit_mpg, newdata = test1,
                             type = "response")

res1 <- table(pred_prob_lda_mpg$class, test1$mpg01)
correct <- res1[1,1]+res1[2,2]
total <- sum(res1)
1-correct/total
```

[1] 0.1025641

The test error obtained using a LDA model is 8.97%.

e-)

```
qda_fit_mpg <- MASS::qda(mpg01~cylinders+weight+displacement
                        +horsepower,
                        data = train1, family = "binomial")

pred_prob_qda_mpg <- predict(qda_fit_mpg, newdata = test1,
                             type = "response")

res_qda <- table(pred_prob_qda_mpg$class, test1$mpg01)
correct_qda <- res_qda[1,1]+res_qda[2,2]
total_qda <- sum(res_qda)
1-correct_qda/total_qda
```

[1] 0.1282051

The test error obtained using a QDA model is 7.69%.

f-)

```
logit_fit_mpg <- glm(mpg01~cylinders+weight+displacement
                    +horsepower,
                    data = train1, family = "binomial")

pred_prob_logit_mpg <- predict(logit_fit_mpg, newdata = test1,
                              type = "response")

pred_prob_logit_mpg[pred_prob_logit_mpg>0.5]<-1
pred_prob_logit_mpg[pred_prob_logit_mpg<0.5]<-0
```

```
res_logit <- table(pred_prob_logit_mpg, test1$mpg01)
correct_logit <- res_logit[1,1]+res_logit[2,2]
total_logit <- sum(res_logit)
1-correct_logit/total_logit
```

[1] 0.1153846

The test error obtained is using the logistic model is 8.97%.

g-)

```
nb_fit_mpg <- e1071::naiveBayes(mpg01~cylinders+weight
                               +displacement+horsepower,
                               data = train1, family = "binomial")

pred_prob_nb_mpg <- predict(nb_fit_mpg, test1)

res_nb <- table(pred_prob_nb_mpg, test1$mpg01)
correct_nb <- res_nb[1,1]+res_nb[2,2]
total_nb <- sum(res_nb)
1-correct_nb/total_nb
```

[1] 0.1025641

The test error obtained is using the naive Bayes model is 6.41%.

h-)

```
tr1_matrix = data.matrix(train1[,c("cylinders", "displacement",
                                   "weight", "horsepower")])
te1_matrix = data.matrix(test1[,c("cylinders", "displacement",
                                   "weight", "horsepower")])

tr1_y = data.matrix(train1$mpg01)
te1_y = data.matrix(test1$mpg01)

predicted_knn_mpg1 <- class::knn(tr1_matrix,
                                test = te1_matrix,
                                cl = tr1_y,
                                k = 1,
                                prob = T)

res_knn1 <- table(predicted_knn_mpg1, te1_y)
correct_knn1 <- res_knn1[1,1]+res_knn1[2,2]
```

```
total_knn1 <- sum(res_knn1)
1-correct_knn1/total_knn1
```

```
[1] 0.1410256
```

The test error obtained is using KNN with K=1 is 11.54%.

```
predicted_knn_mpg5 <- class::knn(tr1_matrix,
                                test = te1_matrix,
                                cl = tr1_y,
                                k = 5,
                                prob = T)

res_knn5 <- table(predicted_knn_mpg5, te1_y)
correct_knn5 <- res_knn5[1,1]+res_knn5[2,2]
total_knn5 <- sum(res_knn5)
1-correct_knn5/total_knn5
```

```
[1] 0.1153846
```

The test error obtained is using KNN with K=5 is 10.27%.

```
predicted_knn_mpg15 <- class::knn(tr1_matrix,
                                  test = te1_matrix,
                                  cl = tr1_y,
                                  k = 15,
                                  prob = T)

res_knn15 <- table(predicted_knn_mpg15, te1_y)
correct_knn15 <- res_knn15[1,1]+res_knn15[2,2]
total_knn15 <- sum(res_knn15)
1-correct_knn15/total_knn15
```

```
[1] 0.1410256
```

The test error obtained is using KNN with K=15 is 12.82%.

```
predicted_knn_mpg10 <- class::knn(tr1_matrix,
                                   test = te1_matrix,
                                   cl = tr1_y,
                                   k = 10,
                                   prob = T)
```



```
res_knn10 <- table(predicted_knn_mpg10, te1_y)
correct_knn10<- res_knn10[1,1]+res_knn10[2,2]
total_knn10 <- sum(res_knn10)
1-correct_knn10/total_knn10
```

```
[1] 0.1153846
```

The test error obtained is using KNN with K=10 is 12.82%.

The best number of nearest neighbors seems to be 5.

Exercises required for MSSC PhD students

1. ISL Sec. 4.8: 4

Solution:

a-) Since we are assuming an uniform distribution, and assuming $x \in [0.05, 0.95]$, then intervals: $[x - 0.05, x + 0.05]$, so length= 0.1. Which means that On average 10% of the observations would be available.

b-) Using $x \in [0.05, 0.95]$, $x1_{length} \times x2_{length} = 0.01$. Therefore, only 1% of the available observations would be used to make a prediction.

c-) Using $p = 100$, gives $0.1^{100} \times 100\%$ of the observations will be available.

d-)

Using the answers from parts a)-c), We notice very quickly that, as dimensionality grows, the likelihood of there being training observations ‘like’ or ‘near’ the test observation X across all p dimensions approaches zero.

Due to the lack of training observations that are “close” across all p dimensions, the K nearest neighbors selected in vast datasets will really not be particularly close. e-) $p = 1; d(length) = 0.1^{1/1} = 0.1$ $p = 2; d(length) = 0.1^{1/2} = 0.32$ $p = 100; d(length) = 0.1^{1/100} = 0.977$

The side length converges to 1 as p increasing, indicating that the hypercube with only 10% of the test observation at its center must be almost the same size as the hypercube with all of the observations. Additionally, it demonstrates that when p rises, observations are concentrated close to the hypercube’s edge and move “further” away from a test observation.

2. ISL Sec. 4.8: 11

Solution:

$$\begin{aligned}
& \log\left(\frac{P_r(Y = k|X = x)}{P_r(Y = K|X = x)}\right) = \log\left(\frac{\pi_k \delta_k(x)}{\pi_K \delta_K(x)}\right) \\
& = \log\left(\frac{\pi_k}{\pi_K}\right) + \pi_k \left(\frac{-1}{2} \log(\Sigma_k) - \frac{-1}{2} (x - \mu_k)^T \Sigma_k^{-1} (x - \mu_k) + \log(\pi_k)\right) \\
& \quad + \pi_K \left(\frac{-1}{2} \log(\Sigma_K) - \frac{-1}{2} (x - \mu_K)^T \Sigma_K^{-1} (x - \mu_K) + \log(\pi_K)\right)
\end{aligned}$$

Then we can separate into constant, linear and quadratic terms

$$\begin{aligned}
& \Rightarrow \log\left(\frac{\pi_k}{\pi_K}\right) - \frac{\pi_k}{2} \log(\Sigma_k) + \pi_k \log(\pi_k) + \frac{\pi_K}{2} \log(\Sigma_K) - \pi_K \log(\pi_K) - \frac{\pi_k}{2} \mu_k^T \Sigma_k^{-1} \mu_k - \frac{\pi_K}{2} \mu_K^T \Sigma_K^{-1} \mu_K \\
& \quad + \pi_k x^T \Sigma_k^{-1} \mu_k + \pi_K x^T \Sigma_K^{-1} \mu_K \\
& \quad - \frac{\pi_k}{2} x^T \Sigma_k^{-1} x - \frac{\pi_K}{2} x^T \Sigma_K^{-1} x \\
& \Rightarrow \log\left(\frac{\pi_k}{\pi_K}\right) - \frac{\pi_k}{2} \log(\Sigma_k) + \pi_k \log(\pi_k) + \frac{\pi_K}{2} \log(\Sigma_K) - \pi_K \log(\pi_K) - \frac{\pi_k}{2} \mu_k^T \Sigma_k^{-1} \mu_k - \frac{\pi_K}{2} \mu_K^T \Sigma_K^{-1} \mu_K \\
& \quad + x^T (\pi_k \Sigma_k^{-1} \mu_k + \pi_K \Sigma_K^{-1} \mu_K) + x^T \left(-\frac{\pi_k}{2} \Sigma_k^{-1} - \frac{\pi_K}{2} \Sigma_K^{-1}\right) x \\
& = a_k + x^T b + x^T c x
\end{aligned}$$

Where

$$a = \log\left(\frac{\pi_k}{\pi_K}\right) - \frac{\pi_k}{2} \log(\Sigma_k) + \pi_k \log(\pi_k) + \frac{\pi_K}{2} \log(\Sigma_K) - \pi_K \log(\pi_K) - \frac{\pi_k}{2} \mu_k^T \Sigma_k^{-1} \mu_k - \frac{\pi_K}{2} \mu_K^T \Sigma_K^{-1} \mu_K$$

,

$$b = (\pi_k \Sigma_k^{-1} \mu_k + \pi_K \Sigma_K^{-1} \mu_K)$$

, and

$$c = -\frac{\pi_k}{2} \Sigma_k^{-1} - \frac{\pi_K}{2} \Sigma_K^{-1}$$

3. ISL Sec. 4.8: 16

First, lets start by generating the binary variable.

```


b_df <- Boston
#Add 1 to column if CRIM > median and 0 otherwise
median_crim <- median(Boston$crim)
b_df$scrim01 <- with(ifelse(crim>median_crim, 1, 0), data=Boston)

```

```

pairs(b_df)

```



```
cor(b_df$crim01,b_df)
```

```
      crim      zn      indus      chas      nox      rm      age
[1,] 0.4093955 -0.436151 0.6032602 0.07009677 0.7232348 -0.1563718 0.6139399
      dis      rad      tax      ptratio      lstat      medv crim01
[1,] -0.6163416 0.6197862 0.6087413 0.2535684 0.4532627 -0.2630167      1
```

It seems to be correlated with everything but the variable “chas”

Let’s separate the data into training and testing.

```
b_sample <- caret::createDataPartition(b_df$crim01, p = 0.8, list = F)
b_train <- b_df[index,]
b_test <- b_df[-index,]
```

Now, let’s see how the models fit.

```
lda_fit_boston <- MASS::lda(crim01 ~.-chas-crim,
                           data = b_train, family = "binomial")

pred_prob_lda_boston <- predict(lda_fit_boston, newdata = b_test,
                                type = "response")

res_b_lda <- table(pred_prob_lda_boston$class, b_test$crim01)
correct_b_lda <- res_b_lda[1,1]+res_b_lda[2,2]
total_b_lda <- sum(res_b_lda)
```

```
1-correct_b_lda/total_b_lda
```

```
[1] 0.1197917
```

The test error using LDA is 12%

```
qda_fit_boston <- MASS::qda(crim01 ~.-chas-crim,
                           data = b_train, family = "binomial")

pred_prob_qda_boston <- predict(qda_fit_boston, newdata = b_test,
                              type = "response")

res_b_qda <- table(pred_prob_qda_boston$class, b_test$crim01)
correct_b_qda <- res_b_qda[1,1]+res_b_qda[2,2]
total_b_qda <- sum(res_b_qda)
1-correct_b_qda/total_b_qda
```

```
[1] 0.06770833
```

The test error using QDA is 6.8%, which is better than LDA.

```
logit_fit_boston <- glm(crim01 ~.-chas-crim,
                      data = b_train, family = "binomial")

pred_prob_logit_boston <- predict(logit_fit_boston, newdata = b_test,
                                type = "response")
pred_prob_logit_boston[pred_prob_logit_boston>0.5]<-1
pred_prob_logit_boston[pred_prob_logit_boston<0.5]<-0

res_logit_b <- table(pred_prob_logit_boston, b_test$crim01)
correct_logit_b <- res_logit_b[1,1]+res_logit_b[2,2]
total_logit_b <- sum(res_logit_b)
1-correct_logit_b/total_logit_b
```

```
[1] 0.08854167
```

The test error using logistic regression is 8.9%, which is better than LDA, but worse than QDA.

```
nb_fit_boston <- e1071::naiveBayes(crim01 ~.-chas-crim,
                                  data = b_train, family = "binomial")

pred_prob_nb_boston <- predict(nb_fit_boston,b_test)
```

```
res_nb_boston <- table(pred_prob_nb_boston, b_test$scrim01)
correct_nb_boston <- res_nb_boston[1,1]+res_nb_boston[2,2]
total_nb_boston <- sum(res_nb_boston)
1-correct_nb_boston/total_nb_boston
```

```
[1] 0.1302083
```

The test error using naive Bayes is 13%, which is worst performing model so far.

```
b_train_m = data.matrix(subset(b_train,select=-c(crim,chas)))
b_test_m = data.matrix(subset(b_test,select=-c(crim,chas)))
train_y = data.matrix(b_train[,14])
test_y = data.matrix(b_test[,14])
```

```
predicted_knn_boston1 <- class::knn(b_train_m,
                                   test = b_test_m,
                                   cl = train_y,
                                   k = 1,
                                   prob = T)

res_knn_b1 <- table(predicted_knn_boston1, test_y)
correct_knn_b1 <- res_knn_b1[1,1]+res_knn_b1[2,2]
total_knn_b1 <- sum(res_knn_b1)
1-correct_knn_b1/total_knn_b1
```

```
[1] 0.07291667
```

The test error using KNN with k=1 is 7.3%, which is the best so far.

```
predicted_knn_boston3 <- class::knn(b_train_m,
                                   test = b_test_m,
                                   cl = train_y,
                                   k = 3,
                                   prob = T)

res_knn_b3 <- table(predicted_knn_boston3, test_y)
correct_knn_b3 <- res_knn_b3[1,1]+res_knn_b3[2,2]
total_knn_b3 <- sum(res_knn_b3)
1-correct_knn_b3/total_knn_b3
```

```
[1] 0.06770833
```

The test error using KNN with k=3 is 6.8%, which is better than k=1.

```

predicted_knn_boston5 <- class::knn(b_train_m,
                                   test = b_test_m,
                                   cl = train_y,
                                   k = 5,
                                   prob = T)

res_knn_b5 <- table(predicted_knn_boston5, test_y)
correct_knn_b5 <- res_knn_b5[1,1]+res_knn_b5[2,2]
total_knn_b5<- sum(res_knn_b5)
1-correct_knn_b5/total_knn_b5

```

```
[1] 0.0625
```

The test error using KNN with $k=5$ is 6.2%, which is better than the previous one.

```

predicted_knn_boston10 <- class::knn(b_train_m,
                                    test = b_test_m,
                                    cl = train_y,
                                    k = 10,
                                    prob = T)

res_knn_b10 <- table(predicted_knn_boston10, test_y)
correct_knn_b10 <- res_knn_b10[1,1]+res_knn_b10[2,2]
total_knn_b10<- sum(res_knn_b10)
1-correct_knn_b10/total_knn_b10

```

```
[1] 0.06770833
```

The test error using KNN with $k=10$ is 6.8%, which is a little worse than $k=5$.

In conclusion, it is possible to see that KNN, with $k=5$, is the best models among the tested ones.

4. In 250 words, summarize of what you learned in the deep learning workshop.

During the deep learning workshop, I was able to gain insight into a few key topics within this vast and rapidly growing field. These included computer vision, natural language processing, and hands-on experience with deep learning frameworks.

We started by learning a little bit about computer vision, which is an application of deep learning that has become increasingly relevant in recent years. This topic covered how deep learning algorithms can be trained to detect and classify objects within images or videos. We also discussed some of the challenges associated with computer vision, including the interpretation of complex visual data, and how deep learning can help overcome these obstacles.

We then moved on to natural language processing, another essential application of deep learning. This topic covered how deep learning algorithms can be trained to process and analyze natural language data, such as text or speech. We explored how this technology has been used to develop chatbots, language translation software, and other language-based applications.

Finally, the workshop provided hands-on experience with deep learning frameworks such as TensorFlow, and Keras. Through this practical application of the concepts we learned, we were able to develop a deeper understanding of how deep learning works in practice.

In conclusion, the workshop provided a comprehensive overview of deep learning, including its applications in computer vision and natural language processing, and practical experience with deep learning frameworks. Getting to see the challenges and solutions when working with deep learning was valuable.

No more optional exercises because no one tried any in the previous two homework sets. Reach out to me if you love to study more machine learning problems.